

Python for Data processing

Lecture 1:

Jupyter, Arrays, tensors and computations

Gleb Ivashkevich

whoami

Gleb Ivashkevich

doing deep learning - time series, satellite imagery

PhD in theoretical physics

6 years in academia doing numerical simulations

6 years in data science and machine learning



datarythmics **effimly**

data driven manufacturing efficiency

Why Python?

Python is:

- simple enough
- flexible
- general purpose
- has huge ecosystem for DS and ML

Why Python?

But it's interpreted! **Isn't it slow?**

Short answer is: **No. It's ok.**

Long answer is: **No, it's not slow**, cause all the heavy lifting is done in **C/C++/Fortran** under the hood. Thank you, Python C API!

Syllabus

Main parts of the course are:

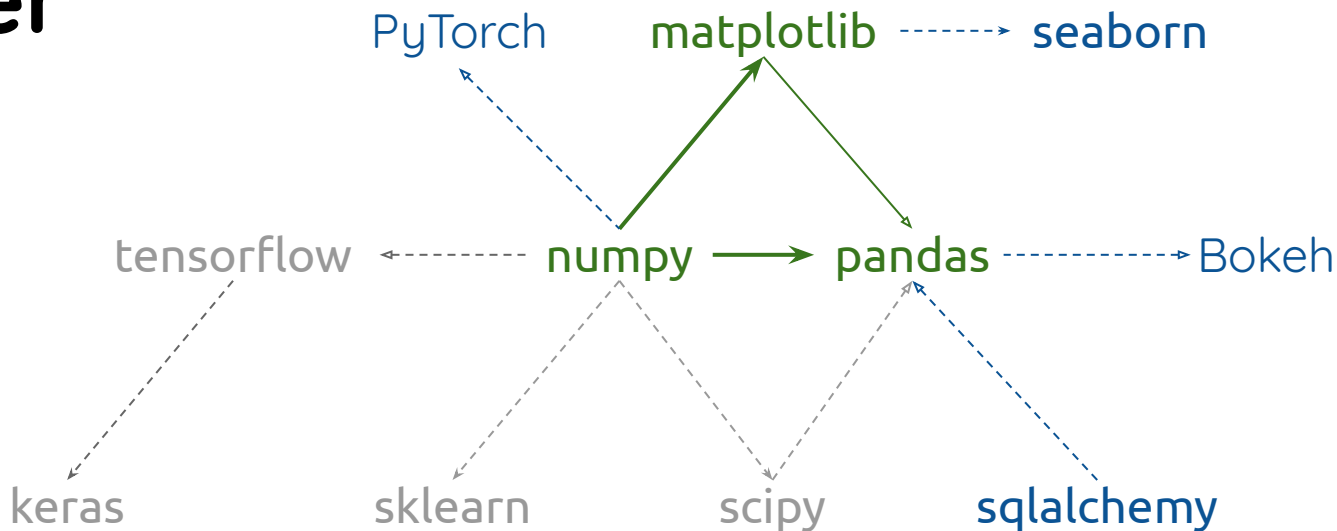
- **numpy** (PyTorch as a topping)
- **matplotlib** (+ seaborn and Bokeh)
- **pandas**
- basics of exploratory data analysis
- tools for reproducibility, project structuring and more

Python ecosystem for ML

Jupyter



papermill



Course logistics

Each week:

- lecture slides (released on Mon, topic based)
- Jupyter notebooks (released on Mon, topic based)
- one graded assignment (released on Wed, week based)
- one or more optional assignments (released during a week)
- + online discussions

Course logistics

Graded assignments:

- we run them with papermill
- they should not fail
- partially autograded
- you have two weeks

Course logistics

Study groups:

- Nov 1: form study groups (if you don't, we assign you randomly)
- Nov 5: tell us, if randomly assigned group is not working for you (logistics, whatever)
- do homework together, discuss, have fun

→let's try it out!

(i.e. we're going to switch to notebook, terminal or whatever)

Resources worth reading

Python for Data Analysis by Wes McKinney

PyData YouTube channel (<https://www.youtube.com/user/PyDataTV>)

From Python to Numpy by Nicolas P. Rougier

<http://www.labri.fr/perso/nrougier/from-python-to-numpy/>

...and there will be more along the way.

Jupyter and other tools

Jupyter

web-based interactive environment

extremely suitable for exploration

originate in IPython project, but is largely **language**

agnostic now

→let's try it out!

Jupyter: a bit of safety

Jupyter server, when running in the cloud/remote machine **should not be open** to the world

Use password and **https** or **ssh** tunneling

numpy:

basics of high performance
arrays

Why numpy?

Pure Python:

- is slow (everything works through Python interpreter)
- lacks strong numerical infrastructure (`math` module? really?)

numpy:

- fast (it's C/C++/Fortran and battle tested BLAS etc. implementations)
- a lot of routines for virtually any generic use case

ndarray

- core data structure in **numpy**
- container with **known number** of elements of the **same** (known) **size**
- supports **indexing** and **vectorized** operations
- allows to **share** data

Creating arrays: naive

- let's use `np.ndarray` directly!
- Or, better, let's create an array from Python sequence

→let's try it out!


Array: basic properties

- shape: `arr.shape`
- type of elements: `arr.dtype`, `arr.itemsize`
- number of dimensions: `arr.ndim`, `ar.size`

Python list vs. ndarray



len = 5
element size = ?



shape = (3, 4) (**elements**)
ndim = 2
element size = fixed

Creating arrays: advanced

We rarely need to create arrays from Python sequences (and `np.ndarray` should be avoided altogether)

Instead we need:

- arrays of **specific structure or type**
- arrays, filled with some **numeric pattern**

→let's try it out!

Array: basic indexing

`numpy` arrays support slicing syntax:

- `a[0, 1]` is ok
- `a[0, :3]` is ok
- `a[1:, :3]` is ok
- `a[0, :-9]` and even this is also ok

→let's try it out!

Array: boolean and fancy indexing

Basic indexing may be (and for large arrays it usually is) insufficient:

- boolean: `a[boolean_mask]`
- fancy: `a[int_idx]` (remember about `np.where`)

→let's try it out!

Array: view vs. copy

Basic indexing returns **view**, fancy and boolean indexing return new **new array**.

But you can mix them.

And it's a bit different for setting values in an array.

→let's try it out!

Array: changing shape

Sometimes we need to change array shape:

- flat vector to row or column vector
- row vector to column vector
- transpose

`arr.reshape`, `np.expand_dims`, `arr.T`

`arr.flatten`, `arr.ravel`

Array: changing type

Sometimes we need to **change array type**:

- to create integer mask
- to reduce memory consumption
- to conform with external API

Array: stack

Sometimes we need to combine multiple arrays:

- to create a matrix from several vectors
- to combine results from different sources

→let's try it out!

Array: universal functions

Fast, vectorized functions, operating element-wise.

- **unary:** `np.sum`, `np.mean` and so on
- **binary:** `np.maximum`, `np.logical_and` and so on

Array: universal functions

ufuncs support common arguments:

- **axis**: operate over this axis
- **where**: masking
- **keepdims**: do not drop reduced dimensions

numpy from inside

ndarray from inside

numpy array
is a container:

```
typedef struct PyArrayObject {  
    PyObject_HEAD  
    char *data; /* Block of memory */  
    PyArray_Descr *descr; /* Data type descriptor */  
    /* Indexing scheme */  
    int nd;  
    npy_intp *dimensions;  
    npy_intp *strides;  
    /* Other stuff */  
    PyObject *base;  
    int flags;  
    PyObject *weakreflist;  
  
} PyArrayObject;
```



ndarray from inside

numpy:

- stores data as a flat chunk of memory
- have indexing scheme on top of that (dimensions and type)
- knows how to step through the memory
- knows the origin

→let's try it out!

ndarray from inside



Consequence #1: cache effects

Data is read from memory in chunks, not element by element



Memory layout may impact performance

→let's try it out!

Consequence #2: copies

Copies are costly



Look at inplace operations

→let's try it out!

Efficient numpy

- use indexing wisely
- avoid loops
- use broadcasting whenever possible
- avoid copies whenever possible
- use inplace operations whenever possible
- **vectorize**

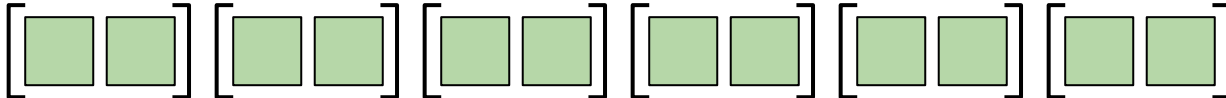
Still slow?

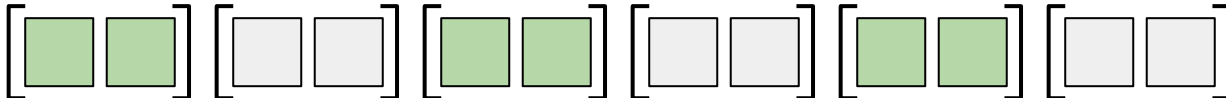
What if the code is so complex, it **gains little** from all the remedies above?

We have tools for that also.

Cython, Numba → optional assignment

View vs. copy


`arrnp.int16` 

`arr[:, :2]` 

`arrnp.int16` **shape** = (6,) **strides** = (2,)
(i) = i * 2

`arr[:, :2]` **shape** = (3,) **strides** = (4,) `arr[:, :2]`
(i) = i * 4 is a **view**

View vs. copy



`arrnp.int16`

shape = (6,) **strides** = (2,)
(i) = i * 2

`arr[:, :2]`

shape = (3,) **strides** = (2,)
(i) = i * 2

`arr[:, :2]`
is a **copy**

Broadcasting

What if input arrays have **different shapes**?

- should we reshape them to common shape before applying some **ufunc**? **No.**
- if possible, ufunc adds missing dimensions and loop through them with stride=0

→let's try it out!

numpy:

basics of linear algebra

Linear algebra: basics

Linear algebra works on vectors (1D), matrices (2D) and tensors (>3D).

Typical operations:

- dot-product of vector and matrix
- matrix operations: invert, get eigenvalues
- decompositions

Linear algebra: `np.linalg`

Entry point for linear algebra operations:

- `np.linalg.inv`, `np.linalg.det`, `np.linalg.trace`
- `np.linalg.eig`
- matrix decompositions

→let's try it out!

Linear algebra: eigenvalues and eigenvectors

The simplest possible decomposition for square matrices

Plays huge role in many algorithms (although, in extended ways)

→let's try it out!

Reading and writing data with **numpy**

Reading and writing numpy array

Array can be saved to a file with `np.save`:

- binary format
- read with `np.load`

→let's try it out!

Reading and writing numpy arrays

Multiple array can be saved to a single file with **np.savez**:

- it's zip, but uncompressed
- use **np.load** to read it (return dict-like object, no data is actually read)

→let's try it out!

Reading and writing text files

`np.loadtxt` and `np.savetxt` are used to read text files (mostly CSV)

But `pandas` is much better in this!

→let's try it out!

Other formats and options

Natively or through `scipy.io`:

- binary data (from files)
- `mat` files
- wav files

Using 3rd party packages:

- HDF5 (`h5py`)
- images (`skimage`, `opencv`)

What we've learned

- creating and indexing arrays
- changing array properties
- calculate with arrays
- basics of linear algebra operations
- I/O

Assignment

- exploring `numpy`: array creation, indexing
- broadcasting
- `ufuncs`

questions?